

DAY 3 ASSIGNMENT

SUMMARY

- HARSHINI V

1. Order Of Execution:

The SQL order of execution is the sequence in which the database management system processes different clauses within a query. This order is essential to ensure that the query produces accurate results and performs efficiently. Typically, the execution starts with the FROM clause, where the database identifies and retrieves the relevant data sources or tables. Following this, the WHERE clause is applied to filter records based on conditions defined by the user, narrowing down the data set to only those entries that meet the specified criteria.

Once the data is filtered, the GROUP BY clause, if present, organizes the data into groups based on certain columns, allowing for aggregation operations like SUM, AVG, or COUNT within each group. After grouping, the HAVING clause is used to filter these groups further, based on specific aggregate conditions. The SELECT clause then defines which columns should appear in the final output, tailoring the result set to meet the user's needs.

Next, the ORDER BY clause is used to sort the results according to specified columns, either in ascending or descending order. Finally, the query execution ends with the LIMIT or OFFSET clause, which restricts the number of rows returned in the result set.

ORDER	CLAUSE	FUNCTION
1	from	Choose and join tables to get base data.
2	where	Filters the base data.
3	group by	Aggregates the base data.
4	having	Filters the aggregated data.
5	select	Returns the final data.
6	order by	Sorts the final data.
7	limit	Limits the returned data to a row count.

This systematic approach to SQL order of execution not only improves resource utilization by reducing the data processed at each stage but also simplifies complex queries, making them easier to read and optimize. Key benefits of understanding this order include:

- Improved accuracy and relevance in query results
- Enhanced performance through query optimization
- More efficient resource usage, reducing server load

2. Sub Total:

Subtotals in SQL queries are crucial in generating reports for fields like sales and finance. A subtotal shows the sum of related data but doesn't represent a grand total. To calculate subtotals, SQL provides the ROLLUP extension, which can generate hierarchical subtotal rows based on column parameters. For example, ROLLUP(SalesYear, SalesQuartes) will calculate subtotals by year and quarter, with an additional grand total row. Using these techniques allows more precise and informative reporting by adding detailed subtotal rows for various levels.

Key Points:

1. **Subtotal Definition:** Subtotals represent the sum of similar data groups without indicating the final total.
2. **SQL Methods:** SQL Server provides extensions like ROLLUP and GROUPING SETS for generating subtotals and grand totals.
3. **GROUPING Function:** This function helps distinguish aggregated columns, allowing us to label subtotal and grand total rows.
4. **Column-Specific Subtotals:** Methods like ROW_NUMBER() with NEWID() can add subtotals for individual columns.
5. **Alternative Grouping Options:** The GROUPING SETS extension offers flexibility by handling multiple grouping configurations in one query.

The GROUPING function plays an essential role in formatting subtotal rows, particularly when combined with conditional logic, such as the SQL CASE statement. This setup can replace NULL values in subtotal rows with descriptive labels like "Subtotal" or "Grand Total." For specific cases, subtotals for individual columns can be calculated using functions like ROW_NUMBER() and NEWID() within a Common Table Expression (CTE), allowing SQL to group and subtotal rows accurately.

Additionally, the GROUPING SETS extension provides an alternative to ROLLUP by allowing multiple grouping levels in one query. This flexibility lets users, for example, create monthly and quarterly sales reports without repetitive UNION statements. Using both ROLLUP and GROUPING SETS gives SQL developers tools to handle subtotals across multiple dimensions, improving the efficiency and readability of complex SQL queries.

3. Stored Procedure:

Stored procedures are a set of precompiled SQL statements that can be saved and reused, enhancing efficiency and reducing redundancy in database operations. By storing SQL queries as procedures, users can execute complex queries multiple times without rewriting the SQL code. This is especially useful for frequently used queries, such as retrieving data from specific tables or performing routine database updates. Stored procedures can also include parameters, which allow them to dynamically adjust based on input values, making them a flexible and powerful tool in SQL.

Structure and Syntax:

Stored procedures follow a straightforward syntax, beginning with the CREATE PROCEDURE statement, followed by the procedure name and the SQL code to be executed. For example:

```
CREATE PROCEDURE procedure_name AS sql_statement GO;
```

To execute the stored procedure, the EXEC command is used:

```
EXEC procedure_name;
```

This basic structure allows for streamlined execution and simplifies the process of running the same query repeatedly. By saving the procedure once, users avoid redundancy and reduce the potential for errors in their code.

Parameterization for Flexibility:

One of the most significant advantages of stored procedures is their ability to accept parameters. Parameters make stored procedures dynamic, allowing them to perform operations based on specific criteria provided at runtime. For example, a stored procedure might retrieve customer records from a particular city:

```
CREATE PROCEDURE SelectAllCustomers @City nvarchar(30) AS SELECT * FROM Customers WHERE City = @City GO;
```

This stored procedure can be executed for a specific city, such as:

```
EXEC SelectAllCustomers @City = 'London';
```

In this case, the procedure only returns records from London. Stored procedures can also handle multiple parameters, allowing for even more refined queries, such as filtering by both city and postal code.

Benefits of Stored Procedures:

Stored procedures offer several benefits:

- **Efficiency:** Since they are precompiled, stored procedures generally execute faster than standalone SQL queries, especially with repetitive tasks.
- **Reduced Redundancy:** They minimize the need to rewrite SQL code for recurring tasks, making them ideal for reports and data retrieval.
- **Improved Security:** Stored procedures can encapsulate SQL code, reducing the exposure of the database structure and enabling better control over access.
- **Maintenance:** Because procedures centralize SQL code, updates or changes can be applied to a single procedure rather than multiple queries across the application.

In sum, stored procedures are a valuable asset in SQL, enhancing performance, security, and flexibility in managing and querying data. They streamline the execution of complex or repetitive tasks, making database operations more consistent and maintainable.

4. Functions In SQL:

It categorizes functions into several types, each serving distinct data manipulation needs. Let's explore each category in detail:

1. String Functions

- **Purpose:** String functions are specifically used for working with text-based data types, such as char and varchar.
- **Common Functions:**
 - **LEN:** Returns the length of a string (e.g., `LEN('WIDESKILLS')` returns 10).
 - **LOWER and UPPER:** Change the case of a string, converting it to either lowercase or uppercase (e.g., `LOWER('JOHN')` returns 'john').
 - **REPLACE:** Substitutes parts of a string with specified characters (e.g., `REPLACE('country', 'y', 'ies')` changes 'country' to 'countries').
 - **REVERSE:** Reverses the characters in a string (e.g., `REVERSE('PATH')` returns 'HTAP').

- **Use Case:** String functions are often applied to format text data, standardize case, and modify data for presentation or further analysis.

2. Date and Time Functions

- **Purpose:** These functions allow for manipulation of date and time data types, such as extracting date parts (e.g., day, month, year), adding time intervals, and performing date calculations.
- **Common Functions:**
 - **GETDATE:** Returns the current system date and time.
 - **DATEADD:** Adds a specified interval to a date, such as months or days (e.g., DATEADD(mm, 2, '2010-02-03') results in '2010-04-03').
 - **DATEDIFF:** Calculates the difference between two dates, based on a specified date part like years or days.
 - **DAY, MONTH, YEAR:** Extract individual parts of a date (e.g., DAY('2010-03-21') returns 21).
- **Use Case:** These functions are useful in applications that track events over time, allowing users to perform scheduling, aging, and trend analysis.

3. Mathematical Functions

- **Purpose:** Mathematical functions handle numerical calculations and transformations on numeric data, including rounding, trigonometry, and logarithmic operations.
- **Common Functions:**
 - **ABS:** Returns the absolute value of a number, eliminating negative signs (e.g., ABS(-77) returns 77).
 - **SIN, COS, TAN:** Provide trigonometric calculations.
 - **CEILING and FLOOR:** Round numbers to the next higher or lower integer, respectively.
 - **EXP:** Calculates the exponential value for a specified parameter.
- **Use Case:** These functions are valuable in scientific, financial, or engineering applications that require precise calculations.

4. Ranking Functions

- **Purpose:** Ranking functions assign sequence numbers to rows in a result set, often based on specific conditions or order. These are helpful for ordering data and creating ranked lists.

- **Common Functions:**
 - **ROW_NUMBER:** Assigns a unique sequential integer to rows, ordered by a specified column.
 - **RANK and DENSE_RANK:** Rank rows based on values in a column; RANK may leave gaps in numbering if there are ties, whereas DENSE_RANK provides consecutive ranking.
 - **NTILE:** Divides rows into a specified number of groups, assigning a group number to each row.
- **Use Case:** Ranking functions are frequently used for leaderboard generation, percentile calculations, and custom sorting.

5. System Functions

- **Purpose:** These functions query system tables, providing information about the server, databases, users, and more. They are particularly useful for administrative tasks.
- **Common Functions:**
 - **HOST_NAME:** Returns the name of the host machine.
 - **SUSER_ID and USER_ID:** Retrieve security identifiers for users and roles.
 - **DB_NAME:** Returns the name of the database based on a specified database ID.
- **Use Case:** System functions are essential for database maintenance, security management, and auditing.

6. Aggregate Functions

- **Purpose:** Aggregate functions summarize data by calculating statistics across a set of rows. They are often used in conjunction with GROUP BY to segment data.
- **Common Functions:**
 - **AVG:** Calculates the average of values in a column.
 - **COUNT:** Counts the number of rows or distinct values.
 - **MIN and MAX:** Return the minimum and maximum values, respectively.
 - **SUM:** Adds up values within a specified column.

- **Use Case:** Aggregate functions are widely used in reporting and data analysis, enabling users to quickly calculate metrics and key performance indicators.

7. Data Grouping and Aggregation Clauses

- **GROUP BY Clause:** Used to group rows that have the same values in specified columns, often paired with aggregate functions to calculate summary values for each group.
 - **Example:** Calculating the minimum and maximum vacation hours for employees by role using GROUP BY emprole.
- **COMPUTE Clause:** An older clause in SQL that generates summarized data, allowing for row-level aggregates in the result set. It differs from GROUP BY in that it can show individual row results alongside aggregates.
- **COMPUTE BY Clause:** Extends COMPUTE by specifying a column to group data, providing subtotals for groups within a dataset.