

PYTHON CASE STUDY

HARSHINI V

Data Processing with Pandas Case Study

Problem Statement:-

Automate the loan eligibility process (real-time) based on customer detail provided while filling the online application form. These details are Gender, Marital Status, Education, Number of Dependents, Income, Loan Amount, Credit History, and others.

The major aim of this notebook is to predict which of the customers will have their loan approved.

- Loading Data in Pandas DataFrame
- Printing rows of the Data
- Printing the column names of the DataFrame
- Summary of Data Frame
- Descriptive Statistical Measures of a DataFrame
- Missing Data Handling
- Sorting DataFrame values
- Merge Data Frames
- Apply Function
- By using the lambda operator
- Visualizing DataFrame

LOADING DATA IN PANDAS DATA FRAME:

```
import pandas as pd
```

```
df = pd.read_csv('C:\\Users\\harsh\\OneDrive\\Documents\\Hexaware\\Role specific class\\csv files\\LoanData.csv')
```

EXPLANATION:

The code imports the pandas library, a powerful Python library used for data manipulation and analysis. It then uses the read_csv method from pandas to load a CSV (Comma-Separated Values) file named LoanData.csv into a pandas DataFrame, which is a two-dimensional, tabular data structure similar to a spreadsheet or database table.

```
# Load Data into Pandas DataFrame
import pandas as pd
df = pd.read_csv('C:\\Users\\harsh\\OneDrive\\Documents\\Hexaware\\Role specific class\\csv files\\LoanData.csv')
```

PRINTING ROWS OF THE DATA:

```
display(df.head())
```

EXPLANATION:

The `display(df.head())` command is used to view the first few rows of the DataFrame `df` in a clean and visually formatted manner, particularly in Jupyter Notebooks or similar environments that support enhanced outputs. The `head()` method retrieves the first 5 rows of the DataFrame by default, providing a quick preview of the data's structure, including the column names, data types, and some sample values.

```
# Printing rows of the Data
display(df.head())
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360.0	1.0	Urban	Y
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0	Rural	N
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	Urban	Y
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	Urban	Y
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	Urban	Y

```
# Printing the column names of the DataFrame
```

PRINTING THE COLUMN NAMES OF THE DATAFRAME:

```
print(df.columns)
```

EXPLANATION:

The `print(df.columns)` command displays the column names of the DataFrame `df`. Each column in a DataFrame represents a specific feature or attribute of the dataset, and this command helps in understanding the structure of the data by listing all the available columns.

```
# Printing the column names of the DataFrame
print(df.columns)
```

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

SUMMARY OF DATA FRAME:

```
print(df.info())
```

EXPLANATION:

The `print(df.info())` command provides a concise summary of the DataFrame `df`, including details about its structure and content. The output lists the total number of rows and columns, the column names, their respective data types (int64, float64, object, etc.), and the non-null count for each column, which helps identify missing data.

```
# Summary of Data Frame
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Loan_ID               614 non-null   object
 1   Gender                601 non-null   object
 2   Married               611 non-null   object
 3   Dependents            599 non-null   object
 4   Education             614 non-null   object
 5   Self_Employed         582 non-null   object
 6   ApplicantIncome       614 non-null   int64
 7   CoapplicantIncome     614 non-null   float64
 8   LoanAmount            592 non-null   float64
 9   Loan_Amount_Term      600 non-null   float64
10  Credit_History        564 non-null   float64
11  Property_Area         614 non-null   object
12  Loan_Status           614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
None
```

DESCRIPTIVE STATISTICAL MEASURES OF A DATAFRAME:

df.describe()

EXPLANATION:

The df.describe() command generates a summary of descriptive statistical measures for the numerical columns in the DataFrame df. The output includes key statistics such as the count (number of non-missing values), mean (average), standard deviation (spread of data), minimum and maximum values, and the 25th, 50th (median), and 75th percentiles. This summary helps in understanding the distribution and variability of the numerical data, identifying potential outliers, and gaining insights into the central tendency and range of the dataset.

```
# Descriptive Statistical Measures of a DataFrame
df.describe()
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.00000	564.000000
mean	5403.459283	1621.245798	146.412162	342.00000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.00000	0.000000
25%	2877.500000	0.000000	100.000000	360.00000	1.000000
50%	3812.500000	1188.500000	128.000000	360.00000	1.000000
75%	5795.000000	2297.250000	168.000000	360.00000	1.000000
max	81000.000000	41667.000000	700.000000	480.00000	1.000000

MISSING DATA HANDLING:

`df.dropna()`

EXPLANATION:

The `df.dropna()` command is used to handle missing data in the DataFrame `df` by removing any rows that contain NaN (Not a Number) or missing values. When this method is applied, it returns a new DataFrame with the rows containing missing values excluded, leaving only the complete cases. This is a simple and effective way to clean data, especially when the proportion of missing values is small and removing them will not significantly affect the dataset.

```
# Missing Data Handling
df.dropna()
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Sta
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0	Rural	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	Urban	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	Urban	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	Urban	
5	LP001011	Male	Yes	2	Graduate	Yes	5417	4196.0	267.0	360.0	1.0	Urban	
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.0	360.0	1.0	Rural	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.0	180.0	1.0	Rural	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.0	360.0	1.0	Urban	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.0	360.0	1.0	Urban	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.0	360.0	0.0	Semiurban	

480 rows × 13 columns

SORTING DATAFRAME VALUES:

```
sorted_df = df.sort_values(by='ApplicantIncome', ascending=False)
```

```
print("Top 5 rows sorted by ApplicantIncome:")
```

```
print(sorted_df.head())
```

EXPLANATION:

The code sorts the DataFrame `df` based on the values in the `ApplicantIncome` column in descending order using the `sort_values` method. The `by='ApplicantIncome'` parameter specifies the column to sort by, and `ascending=False` ensures the sorting is done in descending order, meaning rows with the highest income values appear first. The sorted DataFrame is stored in the variable `sorted_df`. The `print(sorted_df.head())` statement then displays the first five rows of this sorted DataFrame, showing the applicants with the highest incomes.

```
# Sorting DataFrame values
# Sort the dataset by ApplicantIncome in descending order
sorted_df = df.sort_values(by='ApplicantIncome', ascending=False)
print("Top 5 rows sorted by ApplicantIncome:")
print(sorted_df.head())
```

Top 5 rows sorted by ApplicantIncome:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
409	LP002317	Male	Yes	3+	Graduate	No	
333	LP002101	Male	Yes	0	Graduate	NaN	
171	LP001585	NaN	Yes	3+	Graduate	No	
155	LP001536	Male	Yes	3+	Graduate	No	
185	LP001640	Male	Yes	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
409	81000	0.0	360.0	360.0	
333	63337	0.0	490.0	180.0	
171	51763	0.0	700.0	300.0	
155	39999	0.0	600.0	180.0	
185	39147	4750.0	120.0	360.0	

	Credit_History	Property_Area	Loan_Status
409	0.0	Rural	N
333	1.0	Urban	Y
171	1.0	Urban	Y
155	0.0	Semiurban	Y
185	1.0	Semiurban	Y

MERGE DATA FRAMES:

```
df1 = pd.read_csv('C:\\Users\\harsh\\OneDrive\\Documents\\Hexaware\\Role specific class\\csv files\\LoanData.csv')
```

```
df = pd.merge(df,df1)
```

```
print(df)
```

EXPLANATION:

The code demonstrates how to merge two DataFrames, df and df1, using the pd.merge() function from pandas. First, the df1 DataFrame is loaded from a CSV file, just like df. The pd.merge(df, df1) function combines the two DataFrames based on a common column or index. By default, merge() will perform an inner join, meaning only rows with matching values in the common columns from both DataFrames will be retained.

```
# Merge Data Frames
df1 = pd.read_csv('C:\\Users\\harsh\\OneDrive\\Documents\\Hexaware\\Role specific class\\csv files\\LoanData.csv')
df = pd.merge(df,df1)
print(df)
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	\
0	LP001002	Male	No	0	Graduate	No	
1	LP001003	Male	Yes	1	Graduate	No	
2	LP001005	Male	Yes	0	Graduate	Yes	
3	LP001006	Male	Yes	0	Not Graduate	No	
4	LP001008	Male	No	0	Graduate	No	
...	
609	LP002978	Female	No	0	Graduate	No	
610	LP002979	Male	Yes	3+	Graduate	No	
611	LP002983	Male	Yes	1	Graduate	No	
612	LP002984	Male	Yes	2	Graduate	No	
613	LP002990	Female	No	0	Graduate	Yes	

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	\
0	5849	0.0	NaN	360.0	
1	4583	1508.0	128.0	360.0	
2	3000	0.0	66.0	360.0	
3	2583	2358.0	120.0	360.0	
4	6000	0.0	141.0	360.0	
...	
609	2900	0.0	71.0	360.0	
610	4106	0.0	40.0	180.0	
611	8072	240.0	253.0	360.0	
612	7583	0.0	187.0	360.0	
613	4583	0.0	133.0	360.0	

	Credit_History	Property_Area	Loan_Status
0	1.0	Urban	Y
1	1.0	Rural	N
2	1.0	Urban	Y
3	1.0	Urban	Y
4	1.0	Urban	Y

APPLY FUNCTION:

```
def fun(value):
    if value == 'Graduate':
        return "yes"
    else:
        return "No"
df['Education Status'] = df['Education'].apply(fun)
print(df[['Loan_ID','Gender','Education','ApplicantIncome', 'Education Status']].head())
```

EXPLANATION:

The code defines a custom function fun(value) that takes an input value (in this case, a value from the Education column of the DataFrame). If the value is 'Graduate', the function returns "yes", otherwise, it returns "No". This function is then applied to the Education column of the DataFrame using the apply() method, which applies the function to each element in the column. Finally, the print(df[['Loan_ID','Gender','Education','ApplicantIncome', 'Education Status']].head()) statement displays the first five rows of the DataFrame, showing selected columns including the newly created Education Status column, providing a quick overview of the transformed data.

```
# Apply Function
def fun(value):
    if value == 'Graduate':
        return "yes"
    else:
        return "No"
df['Education Status'] = df['Education'].apply(fun)
print(df[['Loan_ID', 'Gender', 'Education', 'ApplicantIncome', 'Education Status']].head())
```

	Loan_ID	Gender	Education	ApplicantIncome	Education Status
0	LP001002	Male	Graduate	5849	yes
1	LP001003	Male	Graduate	4583	yes
2	LP001005	Male	Graduate	3000	yes
3	LP001006	Male	Not Graduate	2583	No
4	LP001008	Male	Graduate	6000	yes

BY USING THE LAMBDA OPERATOR:

```
df['TotalIncome'] = df.apply(lambda x: x['ApplicantIncome'] + x['CoapplicantIncome'], axis=1)
```

```
print("DataFrame with new column TotalIncome:")
```

```
print(df[['ApplicantIncome', 'CoapplicantIncome', 'TotalIncome']].head())
```

EXPLANATION:

The code uses the lambda function to create a new column, TotalIncome, in the DataFrame df. The lambda $x: x['ApplicantIncome'] + x['CoapplicantIncome']$ expression defines an anonymous function that adds the values of ApplicantIncome and CoapplicantIncome for each row. The apply() method applies this function to each row of the DataFrame (specified by axis=1 for row-wise operation). The result is stored in a new column TotalIncome, which holds the sum of the applicant's and coapplicant's incomes. After adding the new column, the print(df[['ApplicantIncome', 'CoapplicantIncome', 'TotalIncome']].head()) statement displays the first five rows of the DataFrame, showing the original ApplicantIncome and CoapplicantIncome columns alongside the newly created TotalIncome column, which reflects the combined income for each applicant and coapplicant.

```
# By using the lambda operator
df['TotalIncome'] = df.apply(lambda x: x['ApplicantIncome'] + x['CoapplicantIncome'], axis=1)
print("DataFrame with new column TotalIncome:")
print(df[['ApplicantIncome', 'CoapplicantIncome', 'TotalIncome']].head())
```

DataFrame with new column TotalIncome:

	ApplicantIncome	CoapplicantIncome	TotalIncome
0	5849	0.0	5849.0
1	4583	1508.0	6091.0
2	3000	0.0	3000.0
3	2583	2358.0	4941.0
4	6000	0.0	6000.0

VISUALIZING DATAFRAME:

```
import matplotlib.pyplot as plt
```

```
df.plot( x='ApplicantIncome',y='LoanAmount',kind='scatter')
```

EXPLANATION:

The code uses the matplotlib.pyplot library to create a scatter plot that visualizes the relationship between two columns, ApplicantIncome and LoanAmount, from the DataFrame df. The df.plot() method is called with parameters x='ApplicantIncome' and y='LoanAmount' to specify that the ApplicantIncome values should be plotted on the x-axis and the LoanAmount values on the y-axis. The kind='scatter' argument specifies that a scatter plot should be generated, which is useful for visualizing the correlation between two continuous variables. The scatter plot will show individual data points as dots, providing a clear view of how ApplicantIncome relates to the LoanAmount. By visualizing this relationship, you can identify any trends, clusters, or potential outliers in the data.

```
# Visualizing DataFrame
import matplotlib.pyplot as plt
df.plot( x='ApplicantIncome',y='LoanAmount',kind='scatter')
```

```
<Axes: xlabel='ApplicantIncome', ylabel='LoanAmount'>
```

